

Judgment Cannot Be Delegated: A Subject-Preserving Framework for AI Governance

Abstract

Contemporary institutions increasingly treat optimised procedures as decisions in their own right. This article advances an ontological limit claim for AI governance: moral judgment is constitutively personal and therefore non-delegable. Building on a minimal philosophical anthropology—person/thing distinction; irreducibility of *phronēsis*; the person as an end; responsibility as constitutive; and the capacity to initiate—we argue that algorithmic assistance can legitimately expand human deliberation, whereas delegation dissolves the very subject who judges. We situate the claim within current debates on artificial agency and Meaningful Human Control (MHC), and show how a subject-preserving reading supplements tracking/tracing by specifying what must remain human in dignity-touching domains. Two diagnostic cases—criminal-justice risk scoring and AI-steered coverage decisions in healthcare—illustrate a structural tendency to displace judgment by optimisation; they are not offered as empirical proof but as paradigmatic contexts where institutional deference to model outputs risks rendering answerability merely nominal. From our axioms we derive governance theorems for non-delegability, contestability, reversibility, and subsidiarity, and we close with an implication for the institutional formation of those who exercise judgment, which exceeds the scope of this article and is taken up in a companion paper. The result is a framework that welcomes instrumental progress while marking a principled boundary: systems may optimise for us; they may not judge instead of us.

Keywords: AI governance, moral judgment, *phronēsis*, meaningful human control, philosophical anthropology, non-delegability, automated decision-making

1 Introduction

A defendant stands before a court while a proprietary algorithm has already rendered its verdict in the form of a risk score. In 2016 the investigative journalists at ProPublica revealed that COMPAS, a recidivism-prediction instrument deployed across the United States, systematically misclassified Black defendants as high-risk at roughly twice the rate of white defendants ([Angwin et al. 2016](#)). The resulting controversy was instructive less for the bias it exposed—racial disparity in criminal justice is, sadly, no novelty—than for the institutional posture it revealed. Judges who might once have been expected to weigh character, context, and the singular trajectory of a life before them had been furnished with a number, and the number had acquired a gravitational authority of its own. What the COMPAS episode brought to public attention was not merely a technical failure but an institutional habit: the quiet replacement of judgment by procedure.

The pattern extends well beyond criminal justice. Investigations have documented how algorithmic systems used by health insurers to manage care for elderly and disabled patients can systematically deny coverage, overriding physician recommendations and reducing complex clinical assessments to binary outputs. In such cases the algorithm does not assist the decision-maker; it supplants her. The physician’s clinical judgment, informed by years of training and the direct encounter with the patient, is subordinated to a model whose internal logic is neither transparent nor contestable in any meaningful sense. Meanwhile, the broader conceptual landscape has shifted. [Floridi \(2025\)](#) has recently proposed that artificial intelligence be understood as “agency without intelligence”—a form of agency that is multiply realisable and need not presuppose

anything resembling human cognition. This reframing is philosophically productive: it liberates the discussion from sterile debates about machine consciousness and focuses attention on the functional capacities that AI systems actually possess. Yet it also raises a question that Floridi’s framework, by design, leaves open: if agency can be realised in non-intelligent substrates, what is it about *human* agency that makes certain decisions non-delegable?

The most influential attempt to answer this question within the governance literature has come through the concept of Meaningful Human Control (MHC). In their canonical formulation, [Santoni de Sio and Van den Hoven \(2018\)](#) propose that human control over an autonomous system is meaningful when two conditions are met: a *tracking* condition, which requires that the system’s behaviour be responsive to the relevant moral reasons, and a *tracing* condition, which requires that the system’s behaviour be attributable to a responsible human agent. This framework has proven remarkably fruitful, generating a substantial body of work that refines, extends, and critiques its core commitments ([Veluwenkamp 2022](#); [Robbins 2024](#); [Abbink et al. 2024](#)). Nonetheless, ambiguities persist. [Robbins \(2024\)](#) has catalogued the many meanings that “meaningful,” “human,” and “control” can bear in this literature, showing that formal compliance with tracking and tracing conditions does not by itself guarantee that the human element retains the character of genuine judgment rather than ritual endorsement.

This article advances a claim that is at once simple and far-reaching: *moral judgment is constitutively personal and therefore non-delegable(?)*. The claim is ontological, not merely normative. It does not assert that delegation is imprudent, or that current AI systems are insufficiently reliable to be trusted with certain tasks. It asserts that the kind of act we call moral judgment—the situated, responsible, context-sensitive discernment of what ought to be done in a particular case—is bound to the structure of personhood in such a way that to delegate it is not to transfer it but to dissolve it. The dissolution occurs because judgment, unlike information, cannot be separated

from the subject who exercises it: information can be transmitted between systems, but judgment requires the integration of perception, deliberation, and accountability in a single responsible subject. When an institution arranges matters so that a model output functions as the effective decision, what is lost is not accuracy but answerability: the presence of a subject who can say “I judged, and here is why.”

The question, in other words, is not whether AI systems are reliable enough to be trusted with certain decisions; it is whether the kind of act we call moral judgment can survive the transfer from person to system.

The article makes ~~four contributions~~. ~~First,~~ three contributions: (i) it articulates a minimal philosophical anthropology—five axioms drawn from the Western tradition—that grounds the non-delegability claim without relying on any single metaphysical system (Section 3). ~~Second,~~ (ii) it derives from these axioms three governance theorems—non-delegability, subject-preserving design, and subsidiarity—that specify constraints on the permissible architecture of human-AI systems (Section 4). ~~Third,~~ (iii) it offers two diagnostic cases—criminal-justice risk scoring and healthcare coverage—not as empirical proof but as paradigmatic illustrations of the structural tendency to displace judgment by optimisation (Section 5). ~~Fourth, it identifies education as the decisive institutional site where the capacity for judgment is either cultivated or allowed to atrophy (Section 6)~~ The argument closes with an implication for the institutional formation of those who exercise judgment, which exceeds the scope of this article and is developed in a companion paper currently under review.

The argument proceeds as follows. Section 2 surveys the state of the art in artificial agency, Meaningful Human Control, and post-principlist AI ethics, and positions the present contribution within those debates. Section 3 states the axiomatic framework. Section 4 derives the governance theorems. Sections ~~5 and 6 develop~~ develops the diagnostic cases ~~and the educational implications, respectively~~. Section ~~7~~ 6 addresses objections. Section ~~8~~ 7 concludes.

2 State of the Art and Positioning

The philosophical literature on AI governance has grown rapidly over the past decade, but three strands are particularly relevant to the present inquiry: the debate over artificial agency and its relationship to responsibility, the development of Meaningful Human Control as a governance concept, and the movement beyond high-level principle lists toward more substantive normative frameworks. We take each in turn before specifying how the present article relates to them.

2.1 Algorithmic decision-making, artificial agency, and responsibility

The question of whether artificial systems can properly be said to possess agency has been a central preoccupation of philosophy of technology since at least the early work of [Floridi and Sanders \(2004\)](#), who argued that artificial agents could be ascribed a minimal form of moral agency understood in terms of their capacity to act in ways that are morally relevant, even if they lack consciousness, intentionality, or freedom in any robust sense. On this view, a thermostat is not a moral agent, but a sophisticated autonomous system that makes decisions affecting human welfare may be—not because it possesses an inner life, but because the functional profile of its behaviour satisfies the criteria for agency at a suitable level of abstraction. This “levels of abstraction” approach proved enormously influential, enabling a productive research programme that avoids the twin pitfalls of anthropomorphism and dismissive reductionism. [Floridi \(2014\)](#) subsequently argued that the pervasive presence of information and communication technologies constitutes a “fourth revolution” in human self-understanding, one that blurs the boundary between online and offline, and between human and artificial agents, in ways that demand a reconception of the moral landscape.

More recently, building on a sustained inquiry into the nature of artificial agency (see also [Floridi 2023](#)), [Floridi \(2025\)](#) has refined and extended this programme with

the thesis that AI constitutes “agency without intelligence.” The core insight is that agency is multiply realisable: the capacity to act in the world, to process information, and to bring about changes in one’s environment does not require the kind of general intelligence that characterises human cognition. A chess engine, a recommendation algorithm, and a self-driving car all exhibit genuine agency—they are not merely passive tools—but they do so without anything resembling understanding, let alone wisdom. This thesis, published in *Philosophy & Technology* itself, represents a significant advance in conceptual clarity. By decoupling agency from intelligence, Floridi frees the philosophical discussion from unproductive speculation about machine consciousness and redirects attention to the functional and normative questions that matter for governance: what can these systems do, what should they be permitted to do, and who is responsible when they do it?

Yet recognising AI as a form of agency, however analytically valuable, is not sufficient for the governance question that concerns us. [Nyholm \(2018\)](#) has argued persuasively that attributing agency to automated systems, while sometimes descriptively apt, risks obscuring the human initiative that designed, deployed, and maintains those systems. When we speak of “the algorithm deciding,” we may inadvertently create a responsibility gap—a space in which no identifiable human agent bears answerability for the outcome. The political dimensions of this displacement have been explored by [Coeckelbergh \(2020\)](#), who situates AI ethics within broader questions of power, participation, and democratic governance, and by [Coeckelbergh \(2022\)](#), who argues that the deployment of AI systems is never merely technical but always already political, insofar as it reconfigures relationships of authority, visibility, and voice.

The crucial point for our purposes is this: the recognition that AI systems possess a form of agency—genuine, functional, consequential—does not settle the question of what kind of agency is required for moral judgment. One may grant everything that Floridi’s framework establishes about the multiple realisability of agency and

still maintain that the specific form of agency involved in moral judgment—situated, embodied, responsive to the singular other, accountable in the first person—is bound to the structure of personhood in ways that resist functional replication. The present article takes Floridi’s analysis as a starting point, not a terminus. If AI is agency without intelligence, what we need to understand is the specific structure of the intelligence—or, more precisely, the practical wisdom—that moral judgment demands, and why that structure is constitutively tied to personhood.

2.2 Meaningful Human Control

The concept of Meaningful Human Control (MHC) emerged from debates over autonomous weapons systems but has since been generalised to a wide range of AI-governance contexts. In their foundational account, [Santoni de Sio and Van den Hoven \(2018\)](#) articulate two conditions that jointly constitute meaningful human control over an autonomous system. The *tracking condition* requires that the system’s behaviour be reliably responsive to the relevant moral reasons that apply in the situation: if the reasons change, the system’s behaviour should change accordingly. The *tracing condition* requires that the system’s behaviour be attributable to a responsible human agent—that there be someone who can properly be held accountable for what the system does. Together, these conditions are intended to ensure that autonomous systems remain within the sphere of human moral responsibility even when they operate without direct human intervention at the moment of action.

This framework has generated a rich and productive literature. [Veluwenkamp \(2022\)](#) has deepened the analysis of the tracking condition by distinguishing between normative reasons and motivating reasons, arguing that meaningful control requires not merely that the system tracks the right reasons but that the human controller’s endorsement of the system is itself based on the right reasons. In other words, it is not enough that the algorithm happens to produce outputs that align with the relevant

moral considerations; the human who deploys the algorithm must understand why those considerations are relevant and must have endorsed the system’s design on that basis. [Robbins \(2024\)](#) has mapped the multiple ambiguities that attend the concept of MHC, showing that “meaningful,” “human,” and “control” each admit of several interpretations, not all of which are compatible with one another. Which humans must exercise control? What precisely is being controlled—the system’s design, its deployment, or its individual outputs? And when is oversight “meaningful” rather than merely ceremonial? These questions remain substantially open.

The most comprehensive recent treatment is the *Research Handbook on Meaningful Human Control of Artificial Intelligence Systems* ([Abbink et al. 2024](#)), which brings together contributions from philosophy, law, engineering, and cognitive science. The handbook makes clear both the achievements and the limitations of the MHC paradigm: it has provided an indispensable vocabulary for governance debates, but it has not yet resolved the fundamental question of what must remain human—and why—when AI systems are integrated into decision-making processes.

The present article contends that the MHC framework, while necessary, is not sufficient. Tracking and tracing are structural conditions on the relationship between a system and its human controllers, but they do not specify the *character* of the human element that must be preserved. A human operator who rubber-stamps algorithmic outputs may satisfy the tracing condition in a formal sense—the output is attributable to her—without exercising anything that deserves the name of judgment. What is needed, we shall argue, is a subject-preserving reading of MHC: one that specifies not merely that a human must be “in the loop” but that the human must remain the *subject* of the decision, the one who judges rather than merely endorsing.

2.3 Beyond principle lists

The early phase of AI ethics was dominated by the proliferation of high-level principle lists. [Jobin et al. \(2019\)](#) documented 84 such documents, identifying convergence around five principles—transparency, justice and fairness, non-maleficence, responsibility, and privacy—while noting significant divergence in interpretation, scope, and enforcement mechanisms. [Hagendorff \(2020\)](#) subjected these guidelines to critical evaluation and found that the vast majority lacked implementation mechanisms, enforcement procedures, and engagement with the political-economic structures within which AI systems are actually developed and deployed. The result, Hagendorff argued, was a form of “ethics washing”: the appearance of ethical governance without its substance. The political-economic context that these guidelines fail to address has been forcefully articulated by [Zuboff \(2019\)](#), whose analysis of “surveillance capitalism” reveals the structural incentives—the extraction and commodification of behavioural data at industrial scale—that drive the deployment of AI systems in ways that are systematically resistant to governance by voluntary ethical principles alone.

In response to these limitations, a number of more substantive approaches have emerged. [Vallor \(2016\)](#) has developed an account of “technomoral virtue” that draws on Aristotelian ethics to argue that what is needed is not merely a list of principles but the cultivation of character dispositions—practical wisdom, justice, courage, temperance—that enable persons to navigate the novel moral terrain created by emerging technologies. [Dignum \(2019\)](#) has articulated a framework for “Responsible AI” that integrates ethical, legal, and technical considerations into the design process itself, arguing that responsibility must be built into AI systems from the outset rather than imposed as an afterthought. [Stahl \(2021\)](#) has proposed an “ecosystem” approach that situates AI ethics within the broader socio-technical systems in which AI is embedded, attending to institutional, economic, and political factors that shape the development and deployment of AI. And [Mittelstadt et al. \(2016\)](#) have provided a foundational mapping

of the ethical issues raised by algorithms, distinguishing between epistemic, normative, and political concerns and showing how they interact.

These developments represent genuine progress beyond the principlist phase. Yet a residual problem persists, and it is this problem that motivates the present article. Even a system that is transparent, fair, non-maleficent, and designed with responsibility built in may still erode the conditions for moral agency if it is deployed in a manner that displaces human judgment rather than supporting it. A principle-compliant algorithm that effectively makes the decision—because the institutional context ensures that its output is routinely followed without genuine deliberation—satisfies every item on the standard checklist while undermining the very capacity that the checklist is meant to protect. The issue is not compliance but the preservation of the judging subject. These advances clarify the insufficiency of principle-compliance and thereby motivate the turn, developed in Section 3, from governance ethics to the philosophical anthropology of the judging subject.

2.4 Positioning

The present article stands at the intersection of these three strands. From the artificial-agency debate we take the insight that AI systems possess genuine functional agency—[Floridi \(2025\)](#) is right that agency is multiply realisable—but we argue that this insight, far from rendering the human-AI boundary irrelevant, makes it all the more urgent to specify what is distinctive about the human side of the boundary. From the MHC literature we take the tracking/tracing framework as an indispensable structural contribution, but we argue that it requires supplementation by a subject-preserving condition that specifies the character of the human element, not merely its presence. From the post-principlist movement we take the recognition that lists of principles are insufficient, but we go further by grounding our normative claims in a minimal philosophical anthropology—a set of axioms about personhood, practical wisdom,

dignity, responsibility, and the capacity to initiate—rather than in a virtue-ethical, deontological, or consequentialist framework taken in isolation.

~~The axiomatic system on which we draw has been stated and developed in ?.~~ In the present article we state the ~~axioms~~ axiomatic system in condensed form (Section 3), derive ~~their~~ its governance implications (Section 4), and test ~~them~~ it against concrete cases (Sections ~~5 and 6~~). The aim is not to replace existing frameworks but to supplement them with a principled account of the boundary between assistance and delegation—between the use of AI that empowers the judging subject and the use that dissolves her.

3 Axiomatic Framework

The framework rests on five axioms drawn from the Western philosophical tradition. Each axiom identifies a structural feature of personhood that is relevant to the governance question. Together, they constitute a *minimal philosophical anthropology*: the smallest set of commitments about the nature of persons that is required to ground the non-delegability claim. The axioms are stated here in condensed form; ~~their historical and systematic development is presented in ?.~~¹

~~A methodological clarification is in order. The axioms~~ The framework adopts a methodological choice that we wish to make explicit before stating the axioms. The displacement of human judgment by algorithmic optimisation is precisely the situation in which classical distinctions between persons and things, between judgment and computation, between answerability and attribution come most under pressure. To invoke those distinctions in this context is therefore not to assume them but to postulate them; this is the strategy of the present article.

¹A more extended exposition of these axioms, with their philosophical lineages traced separately and at length, has been made publicly available by the author [reference omitted for double-blind review]. The present article does not depend on that exposition: the axiomatic system stands or falls by what is argued here.

~~The five axioms that follow are not premises selected *ad hoc* to produce a desired conclusion. Each articulates a commitment that is independently motivated within the philosophical tradition: personalism (Axiom 1), Aristotelian practical philosophy (Axiom 2), Kantian ethics (Axiom 3), Levinasian phenomenology (Axiom 4), and Arendtian political thought (Axiom 5). What is original in our framework follow from the conjunction of the five axioms and require no further postulate. Both are claims that the article must prove, not assume; we ask the reader to defer judgment on the system until its operations have been seen.~~

The five axioms that follow are not premises selected *ad hoc* to produce a desired conclusion. Each articulates a commitment that is independently motivated within the philosophical tradition: personalism (Axiom 1), Aristotelian practical philosophy (Axiom 2), Kantian ethics (Axiom 3), Levinasian phenomenology (Axiom 4), and Arendtian political thought (Axiom 5). What is original in our framework follow from the conjunction of the five axioms and require no further postulate. Both are claims that the article must prove, not assume; we ask the reader to defer judgment on the system until its operations have been seen.

They are postulates: a minimal set of starting propositions that we propose for acceptance and that we are prepared to defend through the consequences of their negation. The relation of the framework to the philosophical tradition is, in this respect, that of a building to its scaffolding. The tradition has been the genetic cause of the construction; once the structure stands, the scaffolding is no longer load-bearing—and one can imagine alternative scaffoldings that would yield the same structure—because what now sustains the building is the mutual coherence of its parts: the axioms support one another, and it is from this internal coherence, not from the tradition that gave them their names, that the framework draws its strength.

We anticipate two senses in which the system is minimal, both of which we demonstrate at the close of this section, after the axioms have been stated. It is structurally minimal: the negation of any axiom collapses the framework as a whole. And it is sufficient: the governance theorems derived in Section 4), and Arendtian political thought (Axiom 5). What is original in our framework follow from the conjunction of the five axioms and require no further postulate. Both are claims that the article must prove, not assume; we ask the reader to defer judgment on the system until its operations have been seen.

A final clarification of strategy. If the postulated system is more than an arbitrary stipulation, it must produce the distinctions that any adequate account of moral agency in the age of artificial intelligence has to be able to draw—between judgment and computation, between answerability and attribution, between the recognition of someone and the classification of something. If it produces them, the system is not

the residue of inherited categories; it is the structure that makes those categories articulable when the situation casts them into doubt. The remainder of the article proceeds on this hypothesis: that the warrant for the axioms lies, in the end, in what they make it possible to say.

3.1 What the framework must articulate

Before stating the axioms, it is worth specifying what a minimal framework adequate to the question of judgment under algorithmic mediation must be able to articulate. The exercise is not yet axiomatic; it is the formulation of the requirements that any such system must satisfy if it is to do philosophical work in the present situation. Two levels emerge.

The first level concerns the relation of one person to another. A framework adequate to this relation must articulate three things. It must articulate what distinguishes the one who acts, judges, and addresses from that on which action and judgment are exercised—in other words, what marks the categorical difference between a who and a what. It must articulate how a person discerns morally in the particular case: not which rule applies, but how the relevant features of the situation are perceived as morally relevant in the first place. And it must articulate what is owed to another person in virtue of her being a person—the substantive content of the recognition without which moral life is not moral.

The second level concerns what becomes of these requirements when something ~~that is not the individual commitments—each has been defended extensively in its own right—but their conjunction into a minimal system and the derivation of governance consequences that are not contained in any single axiom taken alone. Crucially, each axiom is self-standing: it does not require the full metaphysical apparatus of the tradition from which it derives in order to be sustained. The irreducibility of *phronēsis* (Axiom 2) can be affirmed without commitment to Aristotelian hylomorphism; the~~

~~person-as-end (Axiom 3) holds without presupposing transcendental idealism; the primacy of responsibility (Axiom 4) retains its force without adopting Lévinas's wholesale rejection of ontology. What the framework identifies are not fragments extracted from closed systems but *independent convergences*: claims that traditions with incompatible metaphysical commitments nonetheless reach, each for its own reasons. The fact that such diverse traditions converge on the same structural claims about personhood is not a weakness of the framework but its strongest warrant—it suggests that the claims capture something about the structure of moral agency that no single tradition invented. The axioms should therefore be understood not as deductive premises in a formal calculus but as *conditions of possibility* in a broadly transcendental sense—for the distinctions that any adequate governance framework must be able to draw: the distinction between judgment and computation, between a person and a functional profile, between answerability and attribution. If these distinctions are to be sustained, some set of commitments about what persons are and what they can do must be in place; our axioms articulate the minimal such set.~~

person—a system, an institution, an algorithm—is interposed between the person who decides and the person whose life is affected by the decision. Here two further requirements emerge. The framework must articulate what it takes for the deciding person to remain answerable when the operation that produces the decision is no longer wholly hers. And it must articulate what makes it possible for the deciding person to initiate something that is not already contained in the operation of the system—to introduce, into a procedure that would otherwise close on itself, a judgment whose ground is not in the data.

The relation between the two levels is the philosophical hinge of the framework. The requirements of the second level are not derivable from those of the first: no analysis of person-to-person relations alone yields the precise conditions under which a person remains the subject of a decision when a system mediates. Yet the requirements of

the second level are not arbitrary additions either. The contemporary situation has displaced something that was previously implicit. As long as the relation between persons was unmediated, the requirements of the second level were carried silently by those of the first. When systems began to interpose themselves between deciding subjects and affected persons, what had been implicit had to be stated. That is why this framework needs five axioms rather than three.

The five axioms that follow articulate the five requirements identified here. The first three formulate what the person-to-person level demands; the last two formulate what the person-system level adds. The fact that the latter pair are axioms rather than theorems—that they are postulated alongside the first three rather than derived from them—is the structural cost of taking the contemporary situation seriously.

Axiom 1 (Person/Thing Distinction). *There ~~exists an ontological~~ is a categorical distinction between ~~persons—someone~~ and ~~things—that is—something—between a who and a what—that the framework does not~~ ~~reducible—undertake to a difference of degree~~ ~~ground in complexity, information-processing capacity, or functional performance~~any further substantive predicate.*

~~The person/thing distinction is arguably the most fundamental commitment of moral philosophy. Its lineage runs from Boethius (2004), who~~ A1 is regulative in character. It does not specify what the who is; it postulates that the distinction between the who and the what is presupposed by the practices that constitute moral life—addressing, answering, promising, contesting, accusing, forgiving—and that those practices become unintelligible if the distinction is collapsed. To deny A1 is therefore not to discover an empirical fact about persons but to render unintelligible the very practices in which the denial would have to be made. The framework places this minimum at its foundation rather than offering a definition of personhood.

The lineage of the distinction runs through several traditions that approach it from different angles. Aristotle distinguishes the agent who acts and judges from that on

which action is exercised. Descartes, in the *cogito*, finds that the one who doubts is not of the same order as the *res extensae* over which doubt ranges. Boethius (2004) defined the person as *naturae rationalis individua substantia*—an individual substance of a rational nature—through the early-modern and modern traditions to the contemporary work of Spaemann (2006), who nature—and Spaemann (2006), more recently, argues that the distinction between “someone” and “something” is irreducible to any set of functional predicates. Taylor (1989) has shown how this-the distinction is bound up with the very possibility of moral orientation: to be a person is to inhabit a space of questions about the good, and this inhabitation is not a contingent psychological fact but a structural feature of the kind of being a person is. For our purposes, Axiom 1. None of these accounts is adopted here in its full metaphysical commitment. What the framework retains is the structural recognition they share: there is a difference of kind, not of degree, between the one who can be addressed as someone and that which can only be classified as something.

For the purposes of the present article, A1 means that no degree of functional sophistication can transform a system from a thing into a person; the something into someone. The distinction is categorical, not gradual; and it is regulative, not descriptive. What the framework will say about the who—what counts as judging, what counts as being answerable, what counts as initiating—is the work of the axioms that follow, and the article as a whole will return to A1 once those axioms have done their work.

Axiom 2 (Irreducibility of Phronēsis). *Practical wisdom (phronēsis) is irreducible to the application of general rules to particular cases and requires the situated, affectively informed discernment of the morally salient features of a situation.*

Aristotle (2009) distinguishes phronēsis from both *epistēmē* (theoretical knowledge) and *technē* (productive skill), insisting that the practically wise person grasps the particular in a way that cannot be fully captured by universal principles. Dreyfus (1992)

drew on this Aristotelian insight to argue that expert human performance involves a form of embodied, intuitive engagement with the world that resists formalisation in rule-based systems. More recently, [Nussbaum \(2001\)](#) has argued that emotions are not mere accompaniments to judgment but are themselves forms of evaluative perception—ways of apprehending the morally relevant features of a situation that are integral to practical wisdom. Axiom 2 implies that any system whose operation consists in the application of formal rules, statistical regularities, or optimisation functions to data cannot replicate *phronēsis*, however closely its outputs may approximate those of a wise judge in a given domain.

A serious challenge must be acknowledged. [Dreyfus \(1992\)](#)’s critique was directed at rule-based expert systems. Contemporary large language models (LLMs) are not rule-based; they are statistical-connectionist architectures trained on vast corpora of human-generated text, and their capacity to produce contextually sensitive, apparently “wise” outputs is often striking. Does the emergence of LLMs undermine Axiom 2? We argue that it does not, but that the defence requires a more precise formulation than the classical Dreyfusian argument provides.

The key distinction is between what we may call *moral interpolation* and *moral extrapolation*. An LLM can interpolate: within the boundary conditions defined by its training data, it can produce outputs that approximate morally adequate judgments in cases that resemble those already encountered. This capacity is genuine and should not be dismissed. But *phronēsis*, as Aristotle understood it, operates also—and decisively—in the extrapolatory region: in situations where the morally salient features are unprecedented, where existing categories do not apply without transformation, where the right response must be *created* rather than retrieved. It is important to acknowledge that LLMs do exhibit a form of extrapolation in the technical sense: they generate outputs for inputs not present in their training data by extending learned patterns into unvisited regions. But this is *statistical* extrapolation—the projection of

a curve beyond its data points. It is not *moral* extrapolation: the creation of categories adequate to genuinely unprecedented situations, where the right response must be invented, not projected. The AI extends the curve; the person perceives that the curve no longer fits. In the extrapolatory region, the LLM has no moral categories of its own; it has only the statistical residue of moral categories that persons have previously employed.

Moreover—and this is the critical point—the interpolatory capacity of an AI system is *stable only so long as a person supervises it*. That supervision is itself a moral judgment, and it restores human authorship to the decision: the role of the system is to furnish the elements of judgment—data, patterns, alternatives—but the final decision remains with the person, with no displacement in favour of the machine. It is the supervising person who defines the boundary conditions: who recognises when a case falls within the known territory and when it does not, who perceives the morally salient features that the training data may not have captured, who identifies the moment at which the system’s output ceases to be a reliable approximation and becomes a confident confabulation. When the person is removed—when interpolatory moral judgment is automated—the boundary conditions that made it interpolation are themselves removed. What was interpolation under human supervision becomes extrapolation without it, because the human was the reference that kept the system within the known. The decision to delegate even the interpolatory exercise of moral judgment is therefore itself an *extrapolatory* moral decision—an entry into territory that no training data covers—and each subsequent delegation pushes further into uncharted moral space. The door, once opened, does not lead to a stable room; it leads to a corridor of progressive moral displacement, in which the distinction between interpolation and extrapolation erodes with each step. This is why Axiom 2 states not only that AI systems *cannot* replicate phronēsis in the extrapolatory domain but that the delegation of moral judgment even in the interpolatory domain is structurally

unstable: it is a one-way threshold whose crossing transforms the epistemic conditions under which all subsequent judgments are made.

A clarification about the structure of this distinction is now in order, especially in light of the rapid evolution of model capabilities. Two frontiers must be kept apart. The first is empirical: the boundary between simpler and more sophisticated interpolation, between shorter and longer projections within the manifold of patterns the model has learned. This frontier is mobile with scale; larger models extend farther, recombine more, and produce outputs increasingly difficult to distinguish from those a wise judge might give. The second frontier is categorical: it separates any operation conducted with given categories from the *creation* of categories adequate to genuinely novel moral situations. This frontier does not move with scale, because what divides the two sides is not the sophistication of the operation but the place from which it is performed. A non-person agent—regardless of architecture or scale—operates with categories: it receives them from training data, from designers, from previously formed human judgments encoded in the corpus, and within that inheritance it can interpolate, recombine, and project. What it cannot do is produce, in response to a situation that does not fit, the conceptual articulation by which the situation begins to make moral sense at all. That act is not an extension of received categories; it is the originating of new ones. A person can perform it because she stands at the place where categories are formed—the place the judging subject occupies whenever something genuinely unprecedented presents itself.

This is also where Axiom 5 first becomes legible as an axiom. In a world in which the relation between persons is unmediated, the capacity to create new moral categories is carried implicitly by the recognition of the other: persons originate categories in their encounters, and no separate axiom is needed to record this. The interposition of a system reveals what was implicit. A system can be authorised to operate within categories, but the categorical frontier—the boundary between using a category and

creating one—coincides with the boundary between system and person. Axiom 5 is the explicit statement of what becomes visible at that boundary.

Axiom 3 (The Person as an End). *Every person possesses a dignity that prohibits treating her merely as a means to an end, and this dignity generates claims that constrain the permissible scope of optimisation.*

The canonical formulation is Kant’s (1998) Formula of Humanity: “Act so that you treat humanity, whether in your own person or in that of another, always as an end and never merely as a means.” Sen (1999) and Nussbaum (2011) have translated this Kantian commitment into the language of capabilities, arguing that human dignity requires not merely formal rights but the substantive conditions for the exercise of central human capabilities. Sandel (2012) has argued, in a related vein, that certain goods—civic participation, education, health, the administration of justice—resist commodification because subjecting them to market logic corrupts their meaning. Axiom 3 implies that there exist “dignity-touching” domains in which the optimisation of outcomes cannot be the sole or overriding consideration, because what is at stake is not merely efficiency but the treatment of persons as persons.

A challenge deserves candid acknowledgment: what would make a machine *incapable* of treating a person as an end? A system can be programmed to follow protocols that respect dignity—to avoid discrimination, to ensure equitable access, to flag vulnerable cases for human attention. These are genuine and valuable functions. But there is a distinction between following a protocol that respects dignity and *perceiving* the other as a person. The protocol is a rule; the perception is an act of moral recognition—what Lévinas calls the encounter with the face. A machine that applies a dignity-respecting protocol does not thereby treat the person as an end; it processes the case in a manner that is *consistent with* treating the person as an end. The difference is not in the outcome but in the moral structure of the act. An institution that relies exclusively on protocol-compliance—however well-designed—to discharge its obligations toward

persons ~~is operating in what we shall later call the agent production mode: formally correct~~, in this respect, formally correct but ontologically empty. Axiom 3 requires not that the right outcome be produced but that a subject *perceive* the person as a person—and this perception is what grounds the possibility of responding to the singular, the exceptional, the case that the protocol did not foresee.

Axiom 4 (Responsibility as Constitutive). *Responsibility is not an external addition to moral agency but a constitutive feature of personhood: to be a person is to be answerable.*

Lévinas (1969) argues that responsibility for the other is the primordial structure of subjectivity—that the face of the other calls me into being as a moral subject before any act of cognition or volition. Jonas (1984) extends this insight to the technological domain, arguing that the unprecedented power of modern technology demands a new ethics of responsibility in which the continued existence of humanity itself becomes an object of moral obligation. Bauman (1989) offers a dark complement to these accounts: his analysis of the Holocaust reveals how bureaucratic organisation can systematically displace responsibility, enabling moral atrocities precisely by ensuring that no individual agent experiences herself as responsible for the outcome.

The force of Axiom 4 for the AI governance question becomes clear when we distinguish *answerability* from *attribution*. An algorithmic system can be attributed a decision in the causal sense: the output was produced by this model, running on these data, deployed by this institution. The MHC literature’s tracing condition captures this dimension well. But answerability requires something that attribution does not: a subject who can say “I judged, and here is why”—who can give an account of her reasoning, defend it under scrutiny, acknowledge its limitations, and accept its consequences as *her own*. This first-person dimension is not a ceremonial addition to the decision; it is what transforms an output into a judgment. When an institution arranges matters so that no identifiable person occupies this position—so that the

decision is everyone’s in general and no one’s in particular—it has not merely created an accountability gap; it has dissolved the structure of moral personhood at the point where it matters most. Axiom 4 implies that any governance arrangement that makes answerability merely nominal—by interposing an opaque algorithmic process between the decision-maker and the affected person—threatens the constitutive structure of moral personhood.

Axiom 5 (Capacity to Initiate). *Persons possess the capacity to begin something new—to act in ways that are not fully determined by prior causes, social roles, or systemic pressures—and this capacity is the basis of both freedom and political life.*

[Arendt \(1958\)](#) calls this the capacity for “natality”: the ability to interrupt the chain of cause and effect and to introduce something genuinely novel into the world. [Ricoeur \(1992\)](#) develops a complementary account of personal identity as narrative self-constitution, in which the capacity to initiate is bound up with the capacity to give an account of oneself, to say who one is and why one acts as one does.

In the context of AI governance, Axiom 5 does distinctive work that the other axioms cannot perform alone. It is the axiom that grounds the possibility of *contestation*. A person who possesses practical wisdom (Axiom 2) and recognises the other as a person (Axioms 1 and 3) may still lack the capacity—or the institutional conditions—to act on that recognition. Axiom 5 specifies what must be present for the subject to move from perception to action: the capacity to deviate from the algorithmically recommended course, to say “the system recommends X but the situation demands Y,” and to bear the consequences of that departure. Without this capacity, contestability (Theorem 2) remains a formal right with no practical substance—a button that exists but that no one presses because the institutional arrangement has rendered pressing it irrational. An AI system cannot initiate in this sense: it can produce an output that differs from a prior output, but it cannot *refuse* its own recommendation on moral grounds, because refusal presupposes a subject who stands outside the optimisation

function and judges it. This is why Axiom 5 is, together with Axiom 4, the axiom that most clearly resists technological replication: machines can process, recommend, and even “contest” a prior output in the engineering sense, but they cannot initiate in the Arendtian sense—they cannot begin something whose ground is not in the data but in the moral perception of the person who acts.

3.2 Minimality and sufficiency

~~The five axioms are not independent postulates of equal logical standing. Axioms 1–3 are primitive: they state the ontological (person/thing distinction), epistemological (irreducibility of *phronēsis*), and ethical (dignity as end) commitments without which the governance question cannot be formulated. Axioms 4 and 5~~ We close the section by verifying what the methodological paragraphs at the opening of this section anticipated: that the system of five axioms is minimal and sufficient. The verification is not a formal proof—these are postulates, not theorems—but it shows that the framework is closed and stable on its own terms.

The correspondence between the requirements articulated in §3.1 and the axioms stated above is direct. Axiom 1 articulates the categorical distinction between a who and a what. Axiom 2 articulates the situated discernment of the morally salient that is not reducible to rule-application. Axiom 3 articulates what is owed to another person in virtue of being a person. Axiom 4 articulates what it takes for the deciding person to remain answerable when the operation of decision is mediated by a system. Axiom 5 ~~are explications: given that persons are categorically distinct from things and must be treated as ends (Axioms~~ articulates what makes it possible for the person to introduce, into the system’s operation, a judgment whose ground is not in the data. Each requirement of §3.1 is articulated by exactly one axiom. The five-to-five correspondence is exact, and there is no requirement left over and no axiom that does not answer to one.

The system is minimal in the sense that the negation of any axiom collapses the framework. Negate Axiom 1 ~~and~~ and the system has no subject to which the other distinctions can refer; the question of judgment loses its addressee. Negate Axiom 2 and the difference between judgment and computation dissolves; what remains is sophisticated optimisation under a new name. Negate Axiom 3 ~~); responsibility is what makes that commitment operative—its negation would render the conjunction inert. Given the and the~~ person/thing distinction ~~and the irreducibility of practical wisdom (Axioms 1 and 2), the~~ loses its normative consequence; categorical difference becomes a mere taxonomy. Negate Axiom 4 and responsibility decomposes into causal attribution; the question of who can answer for the decision becomes a question of who triggered it. Negate Axiom 5 and practical wisdom remains a disposition without enactment; contestability becomes a button that no one presses. The five negations are not equivalent—each removes a different feature—but each leaves the framework unable to articulate at least one of the requirements identified in §3.1.

A subtler point bears mention here. The negation of Axioms 4 or 5 does not, by itself, break the coherence of the first three. In a world in which the relation between persons is unmediated, the three primitive axioms suffice to articulate the moral life: ~~responsibility and the~~ capacity to initiate ~~is what prevents phronēsis from remaining a mere disposition with no possibility of enactment. The relation is not formal deduction but philosophical explication;~~ are carried implicitly by the recognition of the other. ~~The explicit need for~~ Axioms 4 and 5 ~~articulate what is already implicit in the primitive triad. The system is minimal—the~~ structural trace of contemporary mediation, as §3.1 anticipated. This means that a philosophical reader who finds the first three axioms uncontroversial and the last two tendentious may be reading the system from a position the situation no longer supports. The five-axiom structure is the framework's diagnosis of its own moment.

The system is sufficient in the sense that ~~no axiom can be removed without collapsing the framework (see ?, for the full development).~~ the governance theorems of Section 4 follow from the conjunction of the five axioms with no further postulate, and a sixth axiom would have no requirement to articulate. Sufficiency, like minimality, is verified by the structure of the demonstration that follows: if the theorems can be derived, the postulates were enough; if they could not, a sixth would have been needed. The remainder of the article is, in this sense, the test the system must pass.

3.3 Moral judgment defined

From these axioms we derive a working definition: *moral judgment is the situated, responsible, phronetic discernment of what ought to be done in a particular case, exercised by a person who is answerable for the decision and capable of initiating a course of action that the situation demands.* This definition is Aristotelian in spirit but draws on all five axioms. Judgment in this sense is not the mere selection of an option from a pre-given menu; it involves the perception of morally salient features (Axiom 2), the recognition of the affected party as a person rather than a case (Axioms 1 and 3), the assumption of responsibility for the decision (Axiom 4), and the capacity to act on one’s discernment even when systemic pressures point in a different direction (Axiom 5).

This definition makes clear why judgment resists delegation. A system may process data, detect statistical regularities, and produce recommendations that closely approximate the outputs of wise human deliberation—and it may do so more quickly, more consistently, and across larger populations. What it cannot do is *judge*, because judgment is constitutively tied to the kind of being who judges: a person, situated in a moral world, answerable to others, capable of saying “I” and meaning it.

3.4 Assistance versus delegation

The framework does not entail a blanket rejection of AI in decision-making. The crucial distinction is between *assistance* and *delegation*. An AI system assists when it

enlarges the information available to the judging subject, presents alternatives, identifies patterns that might otherwise be overlooked, or flags inconsistencies in the decision-maker’s reasoning—while the decision-maker retains the capacity and the institutional authority to weigh these inputs against her own situated understanding and to arrive at a judgment for which she is genuinely answerable. An AI system has been delegated to when the institutional arrangement is such that the model output functions as the effective decision—when the human “in the loop” is there as a formality, and deviation from the algorithmic recommendation requires justification while compliance does not.

The distinction is not always sharp, and it can be eroded gradually. [Parasuraman and Manzey \(2010\)](#) have documented the phenomenon of “automation bias”: the tendency of human operators to defer to automated outputs even when those outputs are incorrect, and to fail to search for information that contradicts the automated recommendation. Automation bias means that the boundary between assistance and delegation is not stable; institutional and cognitive pressures tend to push it in the direction of delegation unless active countermeasures are in place. This is why the governance theorems derived in Section 4 include not only a non-delegability condition but also contestability and reversibility requirements: structural safeguards designed to prevent the gradual erosion of the assistance/delegation boundary.

4 Governance Theorems

From the axiomatic framework we derive three governance theorems.² These are not empirical generalisations but normative claims about the constraints that the structure of personhood places on the permissible architecture of human-AI systems. Each theorem is stated as a claim, followed by a rationale that connects it to the axioms and to the existing literature.

²We use “theorem” in the sense standard in normative political philosophy—a substantive claim derived from stated premises through reasoned argument—not in the formal-mathematical sense of a proposition proved within an axiomatic calculus. The derivation is philosophical: the theorems articulate governance constraints that follow from the axioms given the structure of the concepts involved.

4.1 Theorem 1: Non-Delegability

In any domain where a decision bears constitutively on the dignity of a person—where what is at stake is the treatment of someone as a person rather than a thing—the decision may not be delegated to an algorithmic system. The system may inform, recommend, or flag; it may not decide.

Rationale. A preliminary clarification is needed: when does a decision “bear constitutively on the dignity of a person”? We propose the following criterion: a decision is dignity-touching when its outcome determines whether an individual is treated as a *someone*—a singular person whose circumstances, history, and vulnerability require perception and response—or as a *something*—a case to be classified, a profile to be scored, a data point to be processed. The criterion is not about the subjective importance of the decision to the affected party (which would make virtually every decision dignity-touching) but about the structural character of the act: does it require the recognition of singularity, or can it be performed without loss on the basis of a functional profile? Sentencing, medical coverage, asylum adjudication, educational placement, and child custody are paradigmatic dignity-touching decisions because in each case what is at stake is not merely an outcome but whether the person before the decision-maker is seen as a person at all.

The non-delegability theorem follows directly from Axioms 1–4. If persons are categorically distinct from things (Axiom 1), if practical wisdom is irreducible to rule-application (Axiom 2), if persons must be treated as ends (Axiom 3), and if responsibility is constitutive of personhood (Axiom 4), then a decision that bears on the dignity of a person must be made by a subject who is capable of phronetic discernment and who is genuinely answerable for the outcome. To delegate such a decision to an algorithm is not to transfer it to a more efficient decision-maker; it is to transform the nature of the act from judgment to computation, and in so doing to eliminate the very feature—answerability—that gives the act its moral character.

A clarification is needed to prevent misreading. The non-delegability theorem is not an indictment of procedural impersonality as such. Impersonal procedures serve indispensable functions: they protect against favouritism, constrain arbitrary discretion, and ensure that like cases are treated alike. Weber’s analysis of bureaucratic rationalisation identified these virtues clearly, and they remain real. What the theorem targets is not impersonality as a *complement* to personal judgment—which is legitimate and often necessary—but impersonality as a *substitute* for it. A sentencing guideline that structures the judge’s deliberation is a complement; an algorithm whose output functions as the sentence is a substitute. The decisions may be formally identical; what differs is whether a subject stood behind the decision and can give an account of it. It is the telos of the institutional arrangement—whether it is oriented toward supporting the judging subject or replacing her—that determines whether impersonality serves or undermines the conditions of personhood.

The historical resonance of this claim is worth underscoring. [Arendt \(1963\)](#) famously characterised the “banality of evil” as a failure of judgment: Eichmann’s moral catastrophe lay not in demonic intent but in his inability—or refusal—to think, to judge, to see the particular human being before him rather than the bureaucratic category to which that being had been assigned. [Bauman \(1989\)](#) deepened this analysis by showing how modernity’s characteristic organisational form—the bureaucracy—systematically displaces responsibility by distributing it across so many nodes that no individual agent experiences herself as the author of the outcome. The delegation of judgment to algorithmic systems risks reproducing this structure in a new key: not the clerk who follows orders without thinking, but the operator who follows the model output without judging. The moral topology is the same. In both cases what is lost is the presence of a subject who stands before the affected person and can say: “I have considered your case, and here is my judgment, and I am answerable for it.”

The non-delegability theorem does not prohibit the use of AI in dignity-touching domains. It prohibits a specific institutional arrangement: one in which the algorithmic output functions as the effective decision and the human role is reduced to endorsement. The boundary it marks is between a tool that serves the judging subject and a system that replaces her. This distinction may appear obvious in principle, but as the diagnostic cases in Section 5 will show, it is routinely violated in practice—not through deliberate design but through the cumulative pressure of efficiency, liability management, and institutional convenience.

4.2 Theorem 2: Subject-Preserving Design

Governance of AI systems must be subject-preserving rather than merely function-preserving. It is not sufficient that a system produce outcomes that are functionally equivalent to those of a human decision-maker; the institutional arrangement must preserve the conditions under which a human subject genuinely exercises judgment.

Rationale. The distinction between subject-preserving and function-preserving design follows from the irreducibility of *phronēsis* (Axiom 2) and the constitutive role of responsibility (Axiom 4). A function-preserving approach asks only whether the AI system’s outputs match those of a competent human decision-maker; if they do, the system may replace the human without loss. A subject-preserving approach asks a different question: does the institutional arrangement preserve the conditions under which a human being can exercise genuine judgment—perception of the morally salient, deliberation about the good, assumption of responsibility for the outcome?

This distinction gives substance to what Meaningful Human Control should mean. Floridi (2025) has established that AI constitutes a genuine form of agency, and the MHC literature has established that this agency must remain under human control. But the question of what makes control “meaningful” remains, as Robbins (2024) has shown, deeply contested. We propose that meaningful control is subject-preserving

control: an arrangement in which the human controller is not merely present in the causal chain but is the *subject* of the decision—the one who perceives, deliberates, judges, and is answerable.

This reading generates three operational conditions that must be satisfied for an institutional arrangement to count as subject-preserving. First, *contestability*: the human decision-maker must have the authority and the practical capacity to contest, override, or depart from the algorithmic recommendation without bearing an asymmetric burden of justification. If following the algorithm requires no justification but departing from it requires extensive documentation, the arrangement is not subject-preserving even if a human is nominally “in the loop.” Second, *reversibility*: decisions that affect persons must be revisable in light of new information, changed circumstances, or the recognition that the original decision was wrong. A system that automates decisions at scale and makes reversal administratively costly or practically impossible undermines the conditions for subject-preserving governance. Third, *intelligibility*: the human decision-maker must be able to understand, at least in broad outline, the basis on which the algorithmic recommendation was produced, so that she can assess its relevance to the case at hand. This does not require full technical transparency—the decision-maker need not be able to audit the model’s source code—but it does require that the recommendation be accompanied by reasons that are intelligible to a person exercising practical wisdom.

An objection may be anticipated: could not a sufficiently sophisticated AI system satisfy these three conditions on its own? A machine could, in principle, implement a contestation mechanism, a reversal procedure, and an explanation module. But the conditions are not merely procedural; they are subject-requiring. Contestability without a subject who perceives that the output is wrong *for reasons the metric does not capture* is a mechanism without a user. Reversibility without a subject who recognises that circumstances have changed in morally relevant ways is a procedure that no one invokes.

Intelligibility without a subject who can assess whether the reasons offered are adequate to *this* case is transparency addressed to no one. The three conditions presuppose what a machine cannot supply: a person who stands outside the optimisation function and evaluates it. Moreover, even if a machine were to “contest” its own prior output, the contestation would draw on the same statistical patterns that generated the original recommendation; it would be a recalculation, not a judgment. Genuine contestation requires a ground that is not in the data—the moral perception of a person who sees what the data do not encode.

4.3 Theorem 3: Subsidiarity

Decisions should be made at the lowest level of abstraction at which a subject capable of phronetic discernment can exercise judgment, and algorithmic aggregation should not displace this level unless doing so is strictly necessary and the conditions of Theorems 1 and 2 are satisfied.

Rationale. The subsidiarity theorem draws on Axiom 5 (the capacity to initiate) and on the broader principle, well established in political philosophy and Catholic social doctrine, that decisions should be made as close as possible to the persons they affect. [Sandel \(2012\)](#) has argued that certain goods resist the logic of aggregation and scaling because their value is constituted by the manner in which they are produced and distributed, not merely by the outcomes they generate. A verdict handed down by a judge who has weighed the particulars of a case is not the same good as a verdict generated by an algorithm, even if the two are extensionally equivalent, because the former embodies a relationship of recognition between the judging subject and the judged person that the latter does not.

The subsidiarity principle has deep roots in traditions of political thought that emphasise the importance of local knowledge, participatory governance, and the irreducibility of the particular to the universal. In the context of AI governance, it

functions as a counterweight to the tendency—powerful in both technological and bureaucratic cultures—to centralise decision-making at ever-higher levels of abstraction on grounds of efficiency, consistency, and scalability. There are, of course, legitimate reasons to aggregate: epidemiological surveillance, resource allocation across large populations, and the detection of systemic patterns all require forms of analysis that no individual decision-maker can perform. The subsidiarity theorem does not deny this. What it insists is that aggregation must not be allowed to displace the level at which phronetic judgment is exercised, and that when it does, compensating mechanisms—contestability, reversibility, intelligibility—must be put in place.

The three theorems are mutually reinforcing. Non-delegability (Theorem 1) specifies the principled boundary; subject-preserving design (Theorem 2) specifies the institutional conditions that must be met at that boundary; and subsidiarity (Theorem 3) specifies the level at which judgment should be exercised. Together, they constitute a framework that does not reject the use of AI in governance but insists that AI be deployed in ways that preserve the conditions for moral agency—for the exercise of judgment by persons who are answerable for what they do.

5 Diagnostic Cases

The governance theorems articulated above are normative, not empirical. They specify what *ought* to be the case, not what *is* the case. Nonetheless, the framework gains force to the extent that it illuminates patterns that are already visible in the deployment of AI systems. We offer two cases—criminal-justice risk scoring and healthcare coverage—not as empirical proof of the theorems but as paradigmatic illustrations of the structural tendency to displace judgment by optimisation. In each case, the question is not primarily whether the algorithm is biased (though it may be) but whether the institutional arrangement preserves the conditions for judgment.

5.1 Criminal-justice risk scoring

In 2016, ProPublica’s investigation of the COMPAS recidivism-prediction instrument revealed that the system produced racially disparate outcomes: Black defendants were classified as high-risk at roughly twice the rate of white defendants, while the converse error—classifying high-risk defendants as low-risk—was more common among white defendants ([Angwin et al. 2016](#); [Larson et al. 2016](#)). The investigation provoked a vigorous debate, both in public discourse and in the technical literature. [Chouldechova \(2017\)](#) demonstrated a striking impossibility result: under conditions of unequal base rates across groups, it is mathematically impossible for a prediction instrument to satisfy the most common formal fairness criteria simultaneously. If the recidivism rate differs between two groups, then no single instrument can simultaneously achieve equal false-positive rates, equal false-negative rates, and equal positive predictive values across those groups. This result shows that the choice of fairness criterion is not a technical question but a normative one—a judgment about which errors matter most and who should bear the burden of the instrument’s inevitable imperfections.

[O’Neil \(2016\)](#) situates the COMPAS case within a broader pattern: the proliferation of opaque, unaccountable algorithmic systems that reinforce existing inequalities while claiming the authority of mathematical objectivity. The phrase “weapons of math destruction” captures a structural dynamic in which the very features that make algorithms attractive to institutions—their speed, their consistency, their apparent neutrality—also make them effective instruments for the displacement of judgment.

From the perspective of our framework, however, the deepest problem with criminal-justice risk scoring is not bias, important as that is. The deepest problem is structural: the institutional arrangement in which a judge receives a risk score, attached to a particular defendant, and must decide how much weight to give it. In principle, the judge retains full authority to depart from the score. In practice, the score exerts a gravitational pull. Judges who follow the score are protected if the defendant reoffends

(“the algorithm recommended detention”); judges who depart from the score bear the full weight of institutional accountability if their decision proves wrong. The asymmetry of justification—compliance needs no defence, departure does—means that the score functions, *de facto*, as the decision (Dressel and Farid 2018; Stevenson and Doleac 2021). The judge’s role is reduced from judging to endorsing, and the phronetic dimension of sentencing—the perception of the particular person, the weighing of circumstances, the exercise of mercy tempered by justice—is displaced by a number.

What distinguishes this diagnosis from the fairness-centred analyses that have dominated the COMPAS debate? A fairness analysis asks whether the algorithm’s error rates are equitably distributed across demographic groups; a transparency analysis asks whether the model’s internal logic can be made intelligible to the affected parties. Both are important, but neither asks the question that our framework foregrounds: whether the institutional arrangement preserves a judging subject at all. An algorithm with perfectly equalised error rates and full interpretability would still violate the conditions of personhood if the institutional context reduced the judge’s role to endorsing its output. The novelty of the present diagnosis lies in this shift of register—from the properties of the instrument to the status of the subject who is supposed to use it.

This is a violation of Theorem 1 (non-delegability): a dignity-touching decision—one that determines whether a person will be deprived of liberty—is effectively delegated to a system that cannot exercise practical wisdom. It is a violation of Theorem 2 (subject-preserving design): the institutional arrangement does not preserve the conditions under which the judge genuinely judges, because the asymmetry of justification makes departure from the algorithmic recommendation costly and compliance costless. And it is a violation of Theorem 3 (subsidiarity): the aggregated statistical model displaces the situated, particularistic judgment that sentencing requires.

5.2 Healthcare coverage

A parallel pattern has emerged in the healthcare domain. [Obermeyer et al. \(2019\)](#) documented a particularly consequential instance of algorithmic racial bias: a widely used algorithm for identifying patients who would benefit from intensive care management systematically underestimated the health needs of Black patients. The algorithm used healthcare costs as a proxy for health needs, and because Black patients in the United States historically face barriers to accessing care, their costs tend to be lower at any given level of illness. The result was that Black patients had to be considerably sicker than white patients before the algorithm flagged them for additional care. When the researchers corrected the bias, the percentage of Black patients identified for extra care more than doubled.

Investigative reporting has further revealed patterns in which algorithmic systems used by insurers to manage benefits for elderly and disabled patients have overridden physician recommendations, denying coverage for rehabilitation, home health care, and skilled nursing on the basis of model predictions that bore little relationship to the patient’s actual clinical trajectory ([Ross 2023](#)). In these cases, the algorithm did not merely introduce a new data point into the physician’s deliberation; it supplanted the physician’s judgment. The treating clinician, who had examined the patient, spoken with the family, and formed a clinical impression informed by years of training, found her recommendation overruled by a system whose internal logic she could neither inspect nor contest.

[Grote \(2021\)](#) has argued that trustworthy medical AI requires systems that “know when they don’t know”—that is, systems with calibrated uncertainty that can signal when their predictions are unreliable and defer to human judgment. This is an important desideratum, but our framework suggests that it does not go far enough. The issue is not merely epistemic—whether the system can reliably identify the limits of its competence—but ontological: whether the institutional arrangement preserves the

conditions under which a clinician exercises the kind of judgment that medical practice, at its best, embodies. Medicine, like criminal justice, is a domain in which what is at stake is not merely the optimisation of outcomes but the encounter between two persons—one suffering, one trained to respond to suffering—and the quality of that encounter is constitutive of the good that medicine provides. An algorithm that maximises health outcomes across a population may still violate the dignity of the individual patient if it does so by displacing the clinician’s situated judgment about what *this* patient, in *these* circumstances, requires.

Here again, the diagnostic contribution of the framework is distinct from that of existing approaches. Epistemic analyses of medical AI—including the important work of Grote (2021)—focus on the system’s reliability, calibration, and capacity to recognise its own limitations. Our framework does not dispute the value of these desiderata but identifies a deeper structural issue: the displacement of the clinical subject. When an insurer’s algorithm overrides a physician’s recommendation, what is lost is not merely information—which could in principle be corrected—but the encounter between a clinician who has examined this patient and a patient who is owed that examination. The harm is ontological, not merely epistemic: the patient is no longer seen by a subject who judges but processed by a system that classifies.

Both diagnostic cases illustrate the same structural dynamic: institutional deference to algorithmic outputs erodes the conditions for genuine judgment, not because the algorithms are necessarily wrong but because the institutional arrangement transforms the decision-maker from a subject who judges into a conduit for optimised outputs. The problem is not technical but architectural—and the remedy must be architectural as well.

6 ~~Education as the Decisive Site~~

If judgment cannot be delegated, then the capacity for judgment must be cultivated. This observation directs attention to education as the decisive institutional site where the question of AI governance is ultimately settled—not in the design of algorithms or the drafting of regulations, but in the formation of the persons who will use, oversee, contest, and be affected by those systems.

Nussbaum (2010) has argued, with considerable urgency, that the global trend toward instrumentalising education—measuring its value solely in terms of economic productivity—threatens the capacities on which democratic life depends: critical thinking, imaginative empathy, the ability to reason about the common good. When humanities and liberal arts are marginalised in favour of technical skills that can be immediately monetised, societies produce graduates who are well-equipped to function within systems but poorly equipped to question those systems, to see the persons affected by them, or to exercise the kind of judgment that democratic governance requires. ? develops a complementary argument: the dominance of measurement in education—the reduction of educational quality to quantifiable indicators—produces a paradox in which we become good at measuring what we value least and poor at valuing what we cannot measure. The “good” in good education, Biesta insists, cannot be captured by metrics; it involves the formation of subjects who are capable of judgment, responsibility, and action in the world.

Rosa (2019) offers a framework for understanding what is at stake in these arguments at a deeper level. Rosa distinguishes between *resonance*—a mode of relating to the world in which the subject is both affected by and responsive to what she encounters—and *alienation*—a mode in which the world appears as a collection of objects to be controlled, optimised, or consumed. Education, on Rosa’s account, is essentially a resonance phenomenon: genuine learning occurs when the student is addressed by the subject matter, when the encounter with a text, a problem, or another person transforms her understanding and expands her capacity for

response. An educational system that reduces learning to the acquisition of measurable competencies risks producing alienation rather than resonance—students who can process information efficiently but who have lost the capacity to be genuinely affected by what they encounter.

The relevance of these arguments to AI governance is direct. If the capacity for moral judgment is not cultivated through education, then no governance framework can compensate for its absence. Theorems 1–3 specify the conditions that must be met for judgment to be preserved in the architecture of human-AI systems, but those conditions presuppose the existence of subjects who are capable of exercising judgment in the first place. A judge who has been trained to apply rules but not to exercise *phronēsis*, a physician who has been trained to follow protocols but not to see the patient as a person, a civil servant who has been trained to process cases but not to initiate action—these are not subjects who can be meaningfully “in the loop.” They are already, in a sense, functioning as algorithms.

The distinction between *person-formation* and *agent-production* is therefore central. Person-formation is education understood as the cultivation of the capacities identified in the axioms: the capacity to recognise persons as persons (Axiom 1), to exercise practical wisdom in the face of particularity and uncertainty (Axiom 2), to treat persons as ends (Axiom 3), to assume and bear responsibility (Axiom 4), and to initiate action that is not fully determined by systemic pressures (Axiom 5). Agent-production is training understood as the optimisation of human performance for predetermined functions within existing systems. Han (2017) has diagnosed the endpoint of this tendency: a regime in which power operates not through external coercion but through the self-optimisation of subjects who have internalised the logic of the system, producing what he calls the “achievement subject” who is simultaneously free and exploited. Both person-formation and agent-production are necessary—societies need competent engineers, clinicians, and administrators—but if agent-production displaces

person-formation, the result is a population of Han's achievement-subjects: technically proficient but morally hollowed out, capable of executing tasks but incapable of judgment.

The concept of person-formation has deep roots in the pedagogical tradition. Benner (2012) situates it within the broader framework of *Bildsamkeit*—the educability of the human being as an open, self-determining subject whose formation cannot be reduced to the acquisition of predetermined competencies. Stein (1989) developed, from a phenomenological standpoint, an account of empathy (*Einfühlung*) as a distinctive form of intentional experience—the capacity to perceive the other's experience as the experience of a *someone*—which she regarded as foundational for all moral and pedagogical relationships. In her later metaphysical work, ? situated this account within a broader ontology of the person as a finite being oriented toward meaning, showing that the capacity for empathic perception is not a contingent psychological trait but a structural feature of personal existence. These accounts converge with ?'s concept of education as “beautiful risk”: a process whose outcomes cannot be guaranteed in advance because it involves the formation of free subjects who may surprise, dissent, and act in ways that were not predetermined. In terms of our framework, person-formation cultivates the capacity to judge; it does not prescribe the content of the judgment. This is why the objection that person-formation is merely instrumentalisation of another type misses the mark: an education that cultivates the capacity to question ends (Axiom 5) and to assume responsibility for one's decisions (Axiom 4) does not predetermine the person but equips her with the conditions for self-determination. The risk that the person so formed will not exercise judgment is real—it is the “beautiful risk” of freedom. But the risk that a person formed exclusively as an agent will exercise moral judgment is, by the logic of the formation itself, structurally excluded.

~~We acknowledge that this section offers a condensed treatment of a complex pedagogical debate; a fuller development, engaging with the institutional and curricular dimensions of person-formation across educational systems, is undertaken in a companion article currently submitted for publication.~~

~~What would a person-formation curriculum look like in practice? It would include, at a minimum, sustained engagement with philosophical ethics—not as a survey of theories to be memorised but as a practice of deliberation in which students confront genuine moral dilemmas and are required to defend their judgments with reasons. It would include epistemology and philosophy of science, so that students understand not only what algorithms do but what they cannot do—the limits of formalisation, the irreducibility of practical knowledge, the difference between correlation and understanding. It would include political philosophy, so that students can think critically about the institutional arrangements within which AI systems are deployed and can ask who benefits, who bears the costs, and who has a voice in the design. It would include training in metacognition—the capacity to monitor one’s own thinking, to recognise when one is deferring to an automated recommendation out of convenience rather than conviction, and to resist the pull of automation bias. And it would include structured practices of contestation with AI: exercises in which students are required to interrogate algorithmic outputs, identify their assumptions and limitations, and formulate alternative judgments that draw on forms of knowledge—narrative, affective, contextual—that the algorithm cannot capture.~~

~~Such a curriculum is not utopian; elements of it already exist in many educational programmes. But it requires a commitment that is currently unfashionable: the commitment to education as the formation of persons rather than the production of human capital. In the age of AI, this commitment is not a luxury but a condition of possibility for democratic governance.~~

~~A final clarification: we do not claim that formal education is the only site where persons are formed. The workplace, civic participation, family life, and religious communities all contribute to the formation of judgment. But the educational system occupies a distinctive position: it is where a society determines whether its members will possess the tools to ask the questions that all other formative contexts presuppose. A person who has never been equipped to question ends, to reason about justice, or to perceive the singular behind the general will not spontaneously develop these capacities at the workplace. The educational system does not monopolise formation, but it conditions whether formation will occur elsewhere.~~

6 Objections and Replies

The framework advanced in this article is likely to provoke several objections. We consider five of the most pressing and offer replies that clarify and, where necessary, qualify the argument.

6.1 Objection 1: The alignment objection

If AI systems can be aligned with human values—designed so that their objectives reliably reflect the values and preferences of their human principals—then the non-delegability claim is too strong. A sufficiently well-aligned system would make the same decisions that a wise human would, and delegation would involve no loss.

This objection reflects the aspiration of the AI alignment research programme, which seeks to develop methods for ensuring that AI systems pursue objectives that are beneficial to humans. The aspiration is laudable, and progress in alignment research would be welcome. But the objection rests on a confusion between the content of a decision and the structure of judgment. Even if a system could be designed to produce outputs that are extensionally identical to those of a wise judge—the same sentences, the same coverage decisions, the same grades—it would not follow that delegation

involves “no loss.” The loss is not in the content of the decision but in the *mode of its production*. Judgment, as we have defined it, involves the situated, responsible, phronetic discernment of a person who perceives the morally salient, deliberates about the good, and assumes accountability for the outcome. A system that produces the same output without these features has not replicated judgment; it has simulated it. And the difference matters, because what the affected person is owed—in sentencing, in medical care, in educational assessment—is not merely the right outcome but the encounter with a subject who is answerable for that outcome. Alignment secures content-fidelity; it does not and cannot secure subject-presence. The two are not interchangeable, because the moral quality of a decision in dignity-touching domains is partly constituted by the kind of process through which it is reached: a process in which a person stands behind the decision and can give an account of it.

6.2 Objection 2: The MHC sufficiency objection

The tracking and tracing conditions of Meaningful Human Control already ensure that a responsible human agent is in the loop. Adding a “subject-preserving” condition is redundant.

The MHC framework, as articulated by [Santoni de Sio and Van den Hoven \(2018\)](#), is a significant achievement, and we do not propose to replace it. We propose to supplement it. The tracing condition ensures that the system’s behaviour is *attributable* to a responsible human agent; it does not ensure that the agent has *exercised judgment*. Attribution and judgment are not the same thing. A physician who signs off on an algorithmically generated coverage denial is, in the tracing sense, the responsible agent—the denial is attributable to her. But if the institutional arrangement is such that she never examined the patient, never read the clinical notes, and never exercised clinical judgment about the appropriateness of the denial, then the attribution is formal rather than substantive. She is responsible in the sense that she can be held to account, but

she has not judged. The subject-preserving condition fills a gap in the MHC framework by specifying that the human element must be not merely present but *active*—that the person to whom the decision is traced must have genuinely exercised the capacities that constitute judgment. This is not redundant but complementary: it gives the MHC framework the substantive content it needs to distinguish between genuine oversight and ritual endorsement.

6.3 Objection 3: The human bias objection

Human judgment is riddled with bias—racial bias, gender bias, anchoring effects, availability heuristics. Algorithms, for all their limitations, are at least consistent. Insisting on human judgment in domains like criminal sentencing may simply perpetuate human prejudice under the guise of preserving phronēsis.

This objection has real force, and it should not be dismissed. Human judgment is indeed fallible, and the cognitive science literature has documented extensive and systematic biases in human decision-making. Moreover, in some domains, algorithmic decision-making has been shown to produce outcomes that are, on certain metrics, more equitable than those of human decision-makers. But the objection confuses two distinct questions: the question of aggregate performance and the question of legitimacy. It may be true that, on average, an algorithmic system produces outcomes that are more consistent and, by some measure, fairer than those of individual human judges. But consistency and fairness-by-metric are not the same as justice. Justice in the individual case requires the perception of the particular—the recognition of the person before the court as a singular individual, not a member of a statistical class. And this recognition is, as Axiom 2 insists, irreducible to the application of general rules.

Moreover, the objection underestimates the risk of a different form of bias: automation bias. [Parasuraman and Manzey \(2010\)](#) have shown that human operators routinely defer to automated outputs even when those outputs are incorrect, and that this

tendency is robust across domains and levels of expertise. Replacing one form of bias (human prejudice) with another (automation bias) does not represent progress if the latter is less visible, less contestable, and less amenable to correction. The appropriate response to human bias is not to remove the human but to support the human: to provide training, institutional structures, and accountability mechanisms that mitigate bias while preserving the conditions for judgment. The distinction between aggregate performance and individual legitimacy is crucial: even if an algorithm produces better outcomes on average, the individual defendant is owed a judgment, not a prediction.

6.4 Objection 4: The scalability objection

The non-delegability claim may be philosophically sound but practically impossible. Modern institutions process millions of decisions—benefit determinations, loan applications, content moderation—and there simply are not enough trained human judges to exercise phronetic judgment in every case. Insisting on non-delegation is tantamount to insisting that these decisions not be made at all.

This objection is practically important, and we take it seriously. Our framework does not require that every administrative decision be made by a philosopher-judge exercising Aristotelian practical wisdom. It requires that *dignity-touching* decisions—those that bear constitutively on the treatment of persons as persons—be subject to genuine human judgment. Not every decision is dignity-touching: an algorithm that routes a parcel to the nearest distribution centre, or that adjusts the temperature in a building, does not implicate the dignity of persons in the relevant sense. The axioms and theorems specify a principled boundary, and much routine administrative processing falls outside it.

Nevertheless, the objection points to a real tension, and it is worth naming honestly. In some domains, the sheer volume of decisions creates pressure to automate, and this pressure is not illegitimate—denying benefits to millions of people while they

wait for a human reviewer is itself a dignity violation. The subsidiarity theorem (Theorem 3) addresses this tension directly: it does not prohibit aggregation but insists that aggregation must not displace the level at which judgment is exercised, and that compensating mechanisms (contestability, reversibility, intelligibility) must be in place. The practical implication is that institutions must invest in the human infrastructure—trained personnel, adequate staffing, reasonable caseloads—that makes genuine judgment possible, rather than using algorithmic automation as a substitute for that investment. A concrete illustration may clarify the point. Consider disability-benefit determinations, where the volume of applications creates genuine pressure to automate. A subject-preserving architecture might take the form of a two-tier system: an algorithmic screening layer that processes initial applications and approves straightforward cases, paired with a mandatory human-review stage for every denial, in which a trained assessor examines the applicant’s specific circumstances, has the authority to override the algorithmic recommendation without bearing an asymmetric burden of justification, and provides a written account of her reasoning. Such an arrangement uses automation where it is efficient and unobjectionable—approving clear-cut claims—while ensuring that the dignity-touching decision—the denial of a benefit—remains in the hands of a subject who judges.

[Sandel \(2012\)](#) has argued that certain goods resist the logic of commodification and scaling precisely because their value is constituted by the manner in which they are produced. A verdict is not a widget; a clinical judgment is not a data point; an educational assessment is not a score. Insisting that these goods cannot be produced at industrial scale is not impractical idealism; it is a recognition of what they are. If the volume of dignity-touching decisions exceeds the capacity of the available human infrastructure, the appropriate response is to expand the infrastructure, not to redefine the decisions as non-dignity-touching so that they can be automated.

6.5 Objection 5: The extended cognition objection

The extended mind thesis (Clark and Chalmers 1998) holds that cognitive processes are not confined to the skull but extend into tools and environmental structures. If cognition is already distributed across brain, body, and artefact, then the sharp boundary between “the person who judges” and “the AI system that assists” is an artefact of Cartesian prejudice. Moral judgment may be genuinely extended into AI systems, and what this article calls “delegation” is better understood as cognitive extension.

The extended mind thesis, as formulated by Clark and Chalmers (1998), applies principally to sub-personal cognitive processes: memory retrieval, calculation, spatial navigation. The canonical examples—Otto’s notebook, the smartphone as externalised memory—involve information storage and retrieval, not the kind of engaged, responsible discernment that constitutes moral judgment. The extension is at the level of *processing*, not at the level of *personhood*.

Moral judgment, as we have characterised it, is not a cognitive process that happens to occur inside a person; it is an act of the person as a whole—an act that involves perceiving the morally salient features of a singular situation (A2), treating the other as an end (A3), assuming responsibility for the outcome (A4), and initiating something not deducible from prior conditions (A5). These are not computational operations that can be offloaded onto external scaffolding in the way that arithmetic can be offloaded onto a calculator. They are constitutive of the kind of agency that personhood involves.

The extended cognition programme enriches our understanding of the sub-personal mechanisms that support judgment—and we welcome the insight that persons are not self-enclosed Cartesian subjects. But the thesis that cognitive *processes* extend does not entail that moral *agency* extends. A physician who consults an AI diagnostic tool has extended her cognitive reach; a physician whose recommendation is overridden by an insurer’s algorithm has not extended her agency but lost it. The boundary that matters for our framework is not the boundary between brain and world—which is

indeed porous—but the boundary between judgment and its displacement. Extended cognition crosses the first boundary; delegation crosses the second.

7 Conclusion

This article has advanced an ontological limit claim for AI governance: moral judgment is constitutively personal and therefore non-delegable. The claim rests on a minimal philosophical anthropology—five axioms about personhood, practical wisdom, dignity, responsibility, and the capacity to initiate—and generates three governance theorems that specify the conditions under which AI may legitimately be integrated into decision-making processes.

The Non-Delegability Theorem (Theorem 1) establishes that in dignity-touching domains, the effective decision must be made by a person who exercises phronetic discernment and who is genuinely answerable for the outcome. The Subject-Preserving Design Theorem (Theorem 2) specifies that governance arrangements must preserve not merely the presence of a human in the causal chain but the conditions under which that human genuinely judges—contestability, reversibility, and intelligibility. The Subsidiarity Theorem (Theorem 3) specifies that decisions should be made at the lowest level of abstraction at which situated judgment can be exercised, and that algorithmic aggregation must not displace this level without compensating safeguards.

We have illustrated these theorems through two diagnostic cases—criminal-justice risk scoring and healthcare coverage—in which the structural tendency to displace judgment by optimisation is already visible, ~~and we have argued that education is the decisive site where the capacity for judgment is either cultivated or allowed to atrophy.~~ The framework does not reject the use of AI in governance. It welcomes the instrumental gains that well-designed AI systems can provide: greater speed, consistency, and the detection of patterns that human cognition cannot grasp. But it insists on a boundary, and the boundary is principled rather than pragmatic. The question is not whether

AI systems are reliable enough to be trusted with certain decisions; it is whether the kind of act we call moral judgment can survive the transfer from person to system. Our answer is that it cannot, because judgment is constitutively tied to the structure of personhood—to the perception of the particular, the assumption of responsibility, the capacity to say “I” and to be answerable for what follows.

At the opening of Section 3 we postulated Axiom 1 as a minimal regulative distinction: there is a categorical difference between someone and something, between a who and a what, that the framework did not undertake to ground in any further substantive predicate. The axioms that followed have given that minimum its content. We can now say, as a result of the path the article has traced, what the postulate could not say at the outset: the framework has shown why moral judgment beyond received categories presupposes personhood. The capacity to respond to the unprecedented—to perceive that received categories do not fit and to form, in their place, the articulation that a singular situation requires—is not an additional predicate that the postulate now acquires; it is the structural feature of the kind of being to whom the postulate refers, made visible by the system as a whole. What the framework has produced is not a definition of personhood but a vantage from which the categorical distinction between someone and something can be seen to do the work it had to do.

The implications extend beyond the design of algorithms and the drafting of regulations. If judgment cannot be delegated, then the institutions that form persons capable of ~~judgment—schools, universities, professional training programmes—bear~~ judgment bear a responsibility that ~~is, in~~ the present moment ~~insufficiently recognised.~~ The choice between person-formation and agent-production is not a curricular detail; it is a civilisational decision about the kind of beings we wish to be and the kind of governance we wish to sustain. ~~insufficiently recognises. We do not, however, undertake the full development of this implication here: addressing the dilemma of education in its proper depth lies beyond the scope of an article whose object is the structure of~~

judgment under algorithmic mediation. The institutional, curricular, and pedagogical dimensions of person-formation are taken up in a companion paper currently under review. What the present article establishes is, more modestly, a constraint on any pedagogical project compatible with its conclusions: the formation of subjects capable of judgment presupposes institutional arrangements not wholly reduced to systemic logic. The substantive specification of those arrangements is not undertaken here.

We close with the formulation that has guided this inquiry from the outset. Optimisation can inform action; only persons can judge.

Acknowledgements. [To be included in the non-anonymised version.]

Declarations

- **Funding:** No funding was received for conducting this study.
- **Conflict of interest:** The author declares no competing interests.
- **Ethics approval:** Not applicable.
- **Data availability:** Not applicable. This article is a work of philosophical analysis and does not involve empirical data.

References

- Abbink D, Amoroso D, Cavalcante Siebert L, et al (2024) Introduction to meaningful human control of artificially intelligent systems. In: Research Handbook on Meaningful Human Control of Artificial Intelligence Systems. Edward Elgar Publishing, p 1–11, <https://doi.org/10.4337/9781802204131.00006>
- Angwin J, Larson J, Mattu S, et al (2016) Machine bias: There’s software used across the country to predict future criminals. and it’s biased against blacks. ProPublica URL <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>

- Arendt H (1958) *The Human Condition*. University of Chicago Press, Chicago
- Arendt H (1963) *Eichmann in Jerusalem: A Report on the Banality of Evil*. Viking Press, New York
- Aristotle (2009) *The Nicomachean Ethics*. Oxford University Press, Oxford, translated by David Ross, revised by Lesley Brown
- Bauman Z (1989) *Modernity and the Holocaust*. Polity Press, Cambridge
- ~~Benner D (2012) *Allgemeine Pädagogik: Eine systematisch-problemgeschichtliche Einführung in die Grundstruktur pädagogischen Denkens und Handelns*, 7th edn. Juventa, Weinheim~~
- ~~Biesta G (2010) *Good Education in an Age of Measurement: Ethics, Politics, Democracy*. Paradigm, Boulder~~
- ~~Biesta GJJ (2013) *The Beautiful Risk of Education*. Paradigm/Routledge, Boulder~~
- Boethius (2004) *The Theological Tractates and The Consolation of Philosophy*. Loeb Classical Library, Harvard University Press, Cambridge, MA, contains *Liber de persona et duabus naturis* (c. 512). Translated by H. F. Stewart, E. K. Rand, and S. J. Tester
- Chouldechova A (2017) Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data* 5(2):153–163. <https://doi.org/10.1089/big.2016.0047>
- Clark A, Chalmers DJ (1998) The extended mind. *Analysis* 58(1):7–19. <https://doi.org/10.1093/analys/58.1.7>
- Coeckelbergh M (2020) *AI Ethics*. MIT Press, Cambridge, MA

- Coeckelbergh M (2022) The political philosophy of AI: An introduction. *Philosophy & Technology* 35(4):84. <https://doi.org/10.1007/s13347-022-00574-8>
- Dignum V (2019) *Responsible Artificial Intelligence: How to Develop and Use AI in a Responsible Way*. Springer, Cham, <https://doi.org/10.1007/978-3-030-30371-6>
- Dressel J, Farid H (2018) The accuracy, fairness, and limits of predicting recidivism. *Science Advances* 4(1):eaao5580. <https://doi.org/10.1126/sciadv.aao5580>
- Dreyfus HL (1992) *What Computers Still Can't Do: A Critique of Artificial Reason*. MIT Press, Cambridge, MA
- Floridi L (2014) *The Fourth Revolution: How the Infosphere Is Reshaping Human Reality*. Oxford University Press, Oxford
- Floridi L (2023) Present: AI as a new form of agency, not intelligence. In: *The Ethics of Artificial Intelligence: Principles, Challenges, and Opportunities*. Oxford University Press, Oxford, <https://doi.org/10.1093/oso/9780198883098.003.0002>
- Floridi L (2025) AI as agency without intelligence: On artificial intelligence as a new form of artificial agency and the multiple realisability of agency thesis. *Philosophy & Technology* 38(1):30. <https://doi.org/10.1007/s13347-025-00858-9>
- Floridi L, Sanders JW (2004) On the morality of artificial agents. *Minds and Machines* 14(3):349–379. <https://doi.org/10.1023/B:MIND.0000035461.63578.9d>
- Grote T (2021) Trustworthy medical AI systems need to know when they don't know. *Journal of Medical Ethics* 47. <https://doi.org/10.1136/medethics-2021-107463>
- Hagendorff T (2020) The ethics of AI ethics: An evaluation of guidelines. *Minds and Machines* 30(1):99–120. <https://doi.org/10.1007/s11023-020-09517-8>

~~Han BC (2017) Psychopolitics: Neoliberalism and New Technologies of Power. Verso, London, original: *Psychopolitik*, 2014. Translated by Erik Butler~~

Jobin A, Ienca M, Vayena E (2019) The global landscape of AI ethics guidelines. *Nature Machine Intelligence* 1(9):389–399. <https://doi.org/10.1038/s42256-019-0088-2>

Jonas H (1984) *The Imperative of Responsibility: In Search of an Ethics for the Technological Age*. University of Chicago Press, Chicago, original: *Das Prinzip Verantwortung*, 1979. Translated by Hans Jonas and David Herr

Kant I (1998) *Groundwork of the Metaphysics of Morals*. Cambridge University Press, Cambridge, original: *Grundlegung zur Metaphysik der Sitten*, 1785. Translated by Mary Gregor

Larson J, Mattu S, Kirchner L, et al (2016) How we analyzed the COMPAS recidivism algorithm. ProPublica URL <https://www.propublica.org/article/how-we-analyzed-the-compas-recidivism-algorithm>

Lévinas E (1969) *Totality and Infinity: An Essay on Exteriority*. Duquesne University Press, Pittsburgh, original: *Totalité et infini*, 1961. Translated by Alphonso Lingis

Mittelstadt BD, Allo P, Taddeo M, et al (2016) The ethics of algorithms: Mapping the debate. *Big Data & Society* 3(2):1–21. <https://doi.org/10.1177/2053951716679679>

Nussbaum MC (2001) *Upheavals of Thought: The Intelligence of Emotions*. Cambridge University Press, Cambridge

~~Nussbaum MC (2010) *Not for Profit: Why Democracy Needs the Humanities*. Princeton University Press, Princeton~~

Nussbaum MC (2011) *Creating Capabilities: The Human Development Approach*. Harvard University Press, Cambridge, MA

- Nyholm S (2018) Attributing agency to automated systems: Reflections on human–robot collaborations and responsibility-loci. *Science and Engineering Ethics* 24(4):1201–1219. <https://doi.org/10.1007/s11948-017-9943-x>
- Obermeyer Z, Powers B, Vogeli C, et al (2019) Dissecting racial bias in an algorithm used to manage the health of populations. *Science* 366(6464):447–453. <https://doi.org/10.1126/science.aax2342>
- O’Neil C (2016) *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Crown, New York
- Parasuraman R, Manzey DH (2010) Complacency and bias in human use of automation: An attentional integration. *Human Factors* 52(3):381–410. <https://doi.org/10.1177/0018720810376055>
- Ricoeur P (1992) *Oneself as Another*. University of Chicago Press, Chicago, original: *Soi-même comme un autre*, 1990. Translated by Kathleen Blamey
- Robbins S (2024) The many meanings of meaningful human control. *AI and Ethics* 4(4):1377–1388. <https://doi.org/10.1007/s43681-023-00320-6>
- ~~Rosa H (2019) *Resonance: A Sociology of Our Relationship to the World*. Polity Press, Cambridge, original: *Resonanz*, 2016. Translated by James C. Wagner~~
- Ross C (2023) UnitedHealth faces class action lawsuit over algorithm accused of wrongly denying care to seniors. STAT News URL <https://www.statnews.com/2023/11/14/unitedhealth-class-action-lawsuit-algorithm-medicare-advantage/>
- Sandel MJ (2012) *What Money Can’t Buy: The Moral Limits of Markets*. Farrar, Straus and Giroux, New York
- Sen A (1999) *Development as Freedom*. Knopf, New York

- Santoni de Sio F, Van den Hoven J (2018) Meaningful human control over autonomous systems: A philosophical account. *Frontiers in Robotics and AI* 5:15. <https://doi.org/10.3389/frobt.2018.00015>
- Spaemann R (2006) *Persons: The Difference between ‘Someone’ and ‘Something’*. Oxford University Press, Oxford, original: *Personen*, 1996. Translated by Oliver O’Donovan
- ~~Stein E (1989) *On the Problem of Empathy*. ICS Publications, Washington, DC, original: *Zum Problem der Einfühlung*, 1917. Translated by Waltraut Stein~~
- ~~Stein E (2002) *Finite and Eternal Being: An Attempt at an Ascent to the Meaning of Being*. ICS Publications, Washington, DC, original: *Endliches und ewiges Sein*, 1950. Translated by Kurt F. Reinhardt~~
- Stevenson MT, Doleac JL (2021) Algorithmic risk assessment in the hands of humans. IZA Discussion Paper No. 12853
- Stahl BC (2021) *Artificial Intelligence for a Better Future: An Ecosystem Perspective on the Ethics of AI and Emerging Digital Technologies*. Springer, Cham, <https://doi.org/10.1007/978-3-030-69978-9>
- Taylor C (1989) *Sources of the Self: The Making of the Modern Identity*. Harvard University Press, Cambridge, MA
- ~~[Author] (2026) [Bibliographic entry redacted for double-blind review].~~
- Vallor S (2016) *Technology and the Virtues: A Philosophical Guide to a Future Worth Wanting*. Oxford University Press, Oxford
- Veluwenkamp H (2022) Reasons for meaningful human control. *Ethics and Information Technology* 24(4):51. <https://doi.org/10.1007/s10676-022-09673-8>

Zuboff S (2019) *The Age of Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power*. PublicAffairs, New York