

Supplementary Reflection

Notes shared with the reviewers of PHTE-D-26-00254

Optional reading. Not necessary for the assessment of the revision.

A note on these notes

These pages collect a process of reflection that the reviewers' reports prompted and that exceeds what a response letter can usefully accommodate. The reflection has been substantive enough that it would have been dishonest to compress it into the formal channels of the revision. It is offered here as optional reading: the manuscript and the response letter together stand on their own. What follows is the dialogue the reviewers' engagement has made possible.

The document is organised in five sections. Section A explains the methodological reframing of the axiomatic system as postulation rather than inheritance. Section B describes the productive inversion that this reframing makes possible: rather than defending classical distinctions, the framework re-grounds them. Section C extends the reframing to a dialogue with the philosophical traditions that name the axioms and with the contemporary philosophy of technology. Section D revisits the rewriting of Axiom 1 as a minimal regulative postulate that is enriched only retrospectively. Section E clarifies the structure of the interpolation/extrapolation distinction in Axiom 2 by separating two frontiers, only one of which is mobile with scale.

A. Postulating, not inheriting

The first concern that Reviewer 1's report places at the centre of the article is the relation of the axiomatic framework to the philosophical tradition it draws on. The reviewer's diagnosis is that the framework, in its previous formulation, claimed too much: to extract from each tradition only what suits the argument, leaving aside the metaphysical commitments that gave those positions their meaning, is a procedure the reviewer has reasonably found unpersuasive.

The diagnosis is correct, and the response we have arrived at is a methodological reframing of the system itself rather than a defence of its previous form. The axioms are no longer presented as positions distilled from traditions. They are presented as postulates: a minimal set of starting propositions that we propose for acceptance and that we are prepared to defend through the consequences of their negation. The traditions remain the genetic cause of the postulates and we still judge those traditions correct in their own terms. But the framework no longer rests on them in the manner of an inheritance.

The figure that has helped us articulate this is the relation of a building to its scaffolding. Construction requires scaffolding; once the building stands, the scaffolding can be removed without the building falling. One can imagine alternative scaffoldings that would yield the same structure. What now sustains the building is not the scaffolding but the mutual coherence of its parts: the axioms support one another, and the framework's strength comes from this internal coherence rather than from the tradition that named the axioms.

This shift has consequences for how the framework legitimates itself. It does not legitimate by appeal to authority. It legitimates by showing what its negation costs: each axiom can be removed only by collapsing some part of the framework, and the conjunction of the five is sufficient for what the framework needs to articulate. Section 3.7 of the revised article makes this verification explicit. The reviewer's request for systematic build-up internal to the article is, in this sense, the request that the framework justify itself rather than borrow its justification.

It also has consequences for how the framework relates to internal disagreements within the tradition. Aristotle and Kant are not reconciled; Kant and Levinas are not reconciled. If the

framework depended on adopting these traditions whole, those internal disagreements would contaminate the framework. As postulates, the axioms are insulated from those disputes: the disputes can be conducted in their own terrain, while the framework operates with the structural minima it requires. We say more about the use of the framework as a vantage on those disputes in Section C.

The reviewer’s objection has therefore not been deflected; it has been answered by reformulation. The framework does not claim that traditions converge independently on the same structural insights—a claim the reviewer rightly resisted. It claims, instead, that a minimal axiomatic system, postulated honestly and defended by its operations, is what the contemporary situation calls for, and that the tradition appears in this articulation as the scaffolding of a building it caused but no longer holds up.

B. The productive inversion

The deeper component of Reviewer 1’s objection is what we have come to call, following the reviewer’s formulation, the *problem background*. The reviewer points out that the situation motivating the article—the contemporary displacement of human judgment by algorithmic optimisation—is the situation in which the classical distinctions on which the framework draws are themselves under pressure. To invoke those distinctions in this context is to risk treating as established what the present moment leads us to question. A framework that proceeds as if those distinctions were given cannot be expected to dispel a doubt the situation itself produces.

We took this objection to require a reorientation of the framework’s argumentative direction. The previous formulation moved from tradition to argument: the classical distinctions were assumed, and the article defended their continuing relevance under contemporary pressures. The new formulation reverses the direction. The classical distinctions are not assumed; they are postulated as a minimal axiomatic system. The article then proceeds to show that this postulated system gives rise to the distinctions that any adequate account of moral agency under algorithmic mediation must be able to draw—between judgment and computation, between answerability and attribution, between the recognition of someone and the classification of something.

If the system does this work, it is not the residue of inherited categories. It is the structure that makes those categories articulable when the situation has cast them into doubt. The situation does not refute the distinctions: it forces them to be re-postulated from a position that does not depend on the inheritance. The framework thus becomes, in the contemporary moment, less a defence of the classical distinctions than a re-grounding of them.

This inversion is, we believe, more philosophically robust than the language of independent convergences we used in the previous round. Convergences across traditions, even when they are real, do not by themselves answer whether we can still accept the distinctions those traditions articulate. What the rewritten methodological paragraphs claim is something different: that the postulated system, in its operation across the rest of the article, produces the distinctions that the situation casts into doubt, and that this productive capacity is the framework’s warrant.

We share this here as the philosophical movement that the reviewer’s objection has prompted. The version of the argument that appears in the manuscript is intentionally compressed—a published article reaches readers without access to this letter, and an extended self-criticism in the paper would unbalance the opening. The full dialectical movement, including the recognition that the *problem background* is the reviewer’s contribution to the article’s argument, belongs here.

C. A dialogue with the tradition—and with the present

The reviewer’s objection has also prompted a line of reflection that exceeds the scope of the present article and that we share here because its place is the dialogue with the reviewer rather than the manuscript itself. Confronting the axiomatic system with the traditions that name it—retaining what is structural, exhibiting what is contingent—has emerged as a possible field of study: a way of revisiting the humanist tradition from a vantage that does not assume it whole but nonetheless recognises it. That exploration is neither a continuation of nor a complement to the present article; it is a parallel undertaking that we cannot and should not pursue here. What follows are notes from that prospective field, organised first by the principal traditions that the article has named (C.1–C.3) and then by two live debates within the contemporary philosophy of technology (C.4–C.5).

C.1 Aristotle: *phronēsis* withouthylomorphism

Confirmation. A2 retains the situated discernment of the particular. Aristotle’s diagnosis of what is lost when *technē* replaces *phronēsis* is what the framework needs.

Friction. How much Aristotelian metaphysics—natural teleology,hylomorphism—must be retained for *phronēsis* to hold? MacIntyre (*After Virtue*) has defended *phronēsis* without natural teleology; Nussbaum (*The Fragility of Goodness*) has retained more; Dunne (*Back to the Rough Ground*) thematises explicitly what survives without the *Physica*.

Movement. The framework leaves this debate open. A2 functions with whatever response succeeds, and whatever metaphysical apparatus turns out to be needed for *phronēsis* to be sustained, the framework will accommodate. What the framework requires is the structural minimum that all parties to the debate retain.

C.2 Kant: humanity as end without transcendental idealism

Confirmation. A3 retains the dignity-as-end as a non-contingent core.

Friction. Does the Formula of Humanity survive when it is detached from transcendental idealism? Korsgaard (*Sources of Normativity, Self-Constitution*) has argued that it does; Wood (*Kant’s Ethical Thought*) has qualified the claim; Allison (*Kant’s Theory of Freedom*) has questioned it.

Movement. Again, the framework does not need to resolve the dispute. A3 is satisfied by any position that retains the structural minimum: persons are owed treatment as ends. Whether the warrant for that requirement is transcendental, capabilities-based, or otherwise is a question the framework can leave open without loss.

C.3 Levinas / Arendt: responsibility and initiation, ontology pending

Confirmation. A4 retains responsibility as constitutive of subjectivity (Levinas); A5 retains natality and the capacity to initiate (Arendt).

Friction. Levinas rejects ontology; Arendt presupposes it. The framework requires both. Bernstein (*Radical Evil*) and Critchley (*Infinitely Demanding*) have thematised the tension.

Movement. The framework does not resolve the dispute. It does, however, make the tension articulable with precision: by separating A4 and A5 and requiring both, the framework names a place at which the Levinas/Arendt disagreement matters operationally. A philosophical resolution lies outside the article’s object; that the framework can articulate the question is itself a result it can show.

The frictions identified above are inherited—they belong to traditions whose internal disagreements predate the situation that motivates the present article. But the system equally permits the articulation of frictions that are live, that arise within the contemporary philosophy of technology itself.

C.4 Floridi: agency without intelligence; the judging subject

Confirmation. The framework retains Floridi’s insight (2025) that agency is multiply realisable. A2–A5 articulate exactly what no realisation of agency by non-personal substrate can provide.

Friction. Floridi’s reframing separates agency from intelligence but leaves relatively open the question of what specifically is lost when judgment is displaced from the human side. The framework names that loss: the capacity to remain answerable under mediation (A4) and the capacity to initiate from outside the data (A5). These are not contestations of Floridi’s framework; they are specifications of what its analysis of agency leaves to be said about the human side of the human/AI boundary.

Movement. The framework does not displace Floridi’s account. It complements it. Where Floridi’s framework establishes that AI has agency, ours establishes what kind of agency it does not have.

C.5 Meaningful Human Control: tracking, tracing, subject-preserving

Confirmation. The framework retains the Meaningful Human Control insight that human control over autonomous systems requires explicit structural conditions (Santoni de Sio & van den Hoven, 2018, refined by Veluwenkamp, 2022, mapped by Robbins, 2024, surveyed in Abbink et al., 2024).

Friction. Tracking and tracing are conditions on the system’s behaviour and on the attributability of its outputs. They are conditions on the system. The framework adds a condition on the subject: the deciding person must remain the subject of the decision, not merely its endorsing point. The MHC framework, in its standard form, does not articulate this third condition, although recent contributions move in its direction.

Movement. The framework does not displace MHC. It supplements MHC with a subject-preserving condition that gives the human element its substantive content. The supplementation is acknowledged in Section 4 of the manuscript, in the rationale for Theorem 2.

We do not claim that this dialogue settles the reviewer’s original objection in his terms; it does not, because his terms ask the article to do work that exceeds its scope. We do claim that the framework, taken as a postulated structure rather than as inherited orthodoxy, makes such a dialogue articulable; and that this articulability is itself a result the framework can show, not merely promise. The system permits the structuring of frictions both in the tradition that names it and in the philosophy of the present—provided that the frictions remain outside its object, not as a weakness but as evidence of its capacity to articulate.

D. Axiom 1 in two times

The most delicate revision in the article concerns Axiom 1. Reviewer 2 had observed that A1, in its previous formulation, retained a circular character: the axiom appeared to define personhood in terms that prejudged what the article wanted to defend.

The diagnosis is correct, and the response has required separating three distinct levels of circularity. There is an *apparent* circularity that arises whenever a postulate names the very thing the rest of the argument is concerned with—this is unavoidable and harmless if the postulate is regulative rather than substantive. There is a *performative* circularity that arises when the postulate ascribes to its subject precisely the property that the conclusion of the argument depends on—this is *petitio principii*. And there is a *substantive* circularity that arises when no formulation of the postulate can avoid presupposing what the argument concludes—this is a deeper kind of difficulty that no reformulation can entirely resolve.

The revised formulation of A1 is designed to be of the first kind without being of the second. At the opening of Section 3, A1 is stated as a minimal regulative postulate: there is a categorical distinction between someone and something, between a who and a what, that the framework

does not undertake to ground in any further substantive predicate. The framework places this minimum at its foundation rather than offering a definition of personhood. What sustains A1 is not a property attributed to persons but the unintelligibility of the practices that constitute moral life—addressing, answering, promising, contesting, accusing, forgiving—if the distinction is collapsed. To deny A1 is therefore not to discover an empirical fact about persons but to render unintelligible the very practices in which the denial would have to be made.

We considered, in our internal deliberations, a more substantive formulation: defining the person as the kind of being capable of issuing moral judgment. We discarded this formulation precisely because it would have made the article’s conclusion analytic. If A1 says *person = capable of moral judgment*, then the thesis *delegating judgment depersonifies* is true by the meanings of its terms. The reviewer’s observation that A1 was circular found, in this respect, a stronger version of the difficulty than perhaps was first apparent: the temptation to define is the temptation to make the argument analytic.

A1 returns at the close of the article, but in a form that we hope makes the asymmetry clear. The conclusion does not redefine personhood. It shows—as a result of the operations the article has performed—why moral judgment beyond received categories presupposes personhood. The capacity to respond to the unprecedented is not a predicate that A1 now acquires; it is a structural feature of the kind of being to whom A1 referred all along, made visible by the system as a whole. The article produces a vantage from which the categorical distinction can be seen to do the work it had to do.

This is, in effect, the article’s productive inversion (Section B above) applied to A1 itself. The minimum at the opening is not enriched by stipulation; it is enriched by what the article has shown about it. The circularity the reviewer detected is, by the new architecture, structurally avoided: the postulate at the beginning does not contain the conclusion, and the result at the end is reconstructive rather than definitional.

E. Two frontiers, not one

Reviewer 2’s gentlest comment was also, in our judgment, the one that most needed addressing in the long term. The reviewer welcomed the distinction between *moral interpolation* and *moral extrapolation* but observed that whether it would hold as models scale is a matter of speculation. The implicit warning is that an argument depending on the present capacities of large language models will be vulnerable as those capacities grow.

The clarification we have added to A2 separates what the original distinction conflated. There are two frontiers at issue.

The first frontier is empirical. It runs between simpler interpolation and more sophisticated interpolation, between shorter and longer projections within the manifold of patterns a model has learned. This frontier moves with scale. We do not deny it: larger models extend farther, recombine more, and produce outputs increasingly difficult to distinguish from those a wise judge might give. The reviewer’s caveat about scaling concerns this frontier, and the reviewer is right to caution against arguments that depend on it.

The second frontier is categorical. It runs between any operation conducted with given categories, however sophisticated, and the creation of categories adequate to genuinely novel moral situations. This frontier does not move with scale. What divides the two sides is not the sophistication of the operation but the place from which it is performed. A non-person agent—regardless of architecture, regardless of training data, regardless of scale—operates with categories. It receives them from the corpus, from designers, from previously formed human judgments. Within that inheritance it can interpolate, recombine, and project in ways that exceed any individual instance of training. What it cannot do is produce, in response to a situation that does not fit, the conceptual articulation by which the situation begins to make moral sense at all. That act is

not an extension of received categories; it is the originating of new ones, and it belongs to the position that only a person occupies.

The framework's argument depends on the second frontier, not the first. It is therefore not vulnerable to scaling: each new generation of models, however capable, remains on the operational side of the categorical frontier. The frontier is structural; what determines the side on which a system falls is not its capability but the place from which it operates.

This articulation also makes visible the connection between A2 and A5. In a world in which the relation between persons is unmediated, the capacity to create new moral categories is carried implicitly by the recognition of the other; persons originate categories in their encounters, and no separate axiom is needed to record this. The interposition of a system reveals what was implicit. A system can be authorised to operate within categories, but the categorical frontier—the boundary between using a category and creating one—coincides with the boundary between system and person. A5 is the explicit statement of what becomes visible at that boundary. The reviewer's caveat has therefore prompted a clarification that, beyond addressing scaling, has reinforced the structural unity of the framework.

Closing note

These pages are, as the title notes, optional reading. The manuscript and the response letter together carry the formal weight of the revision; what is gathered here is the surplus that the reviewers' engagement has produced—pieces of philosophical thinking that grew during the dialogue and that we share with the reviewers in recognition of what their reading has made possible.

We are grateful to both reviewers for the depth of their reading. The article that will be evaluated is the article they have helped us write. If, for any reason, this supplementary document is not consulted, no harm is done to the assessment of the manuscript. If it is consulted, we hope the reader will find that the dialogue has been received seriously enough to be returned in kind.